

# ASSIGNMENT 2

---

## Auto-Scaling Web Application

Development Operations  
26670 (2023 - 2024)

Student: Michael McKibbin

ID: 20092733

## Contents

Planning: Application selection and architecture planning, including architecture diagram. AMI creation .....	4
Installation of nvm, node, and app.js .....	5
VPC and load balancer implementation.....	8
Testing: Test traffic and load balancing demonstration.....	26
Monitoring: Own script. ....	29
APPENDIX A: Automated instance setup.....	30
APPENDIX B: monitor.sh - Instance info.....	39
APPENDIX C: memv2.sh - Cloudwatch. ....	40

## Table of Figures

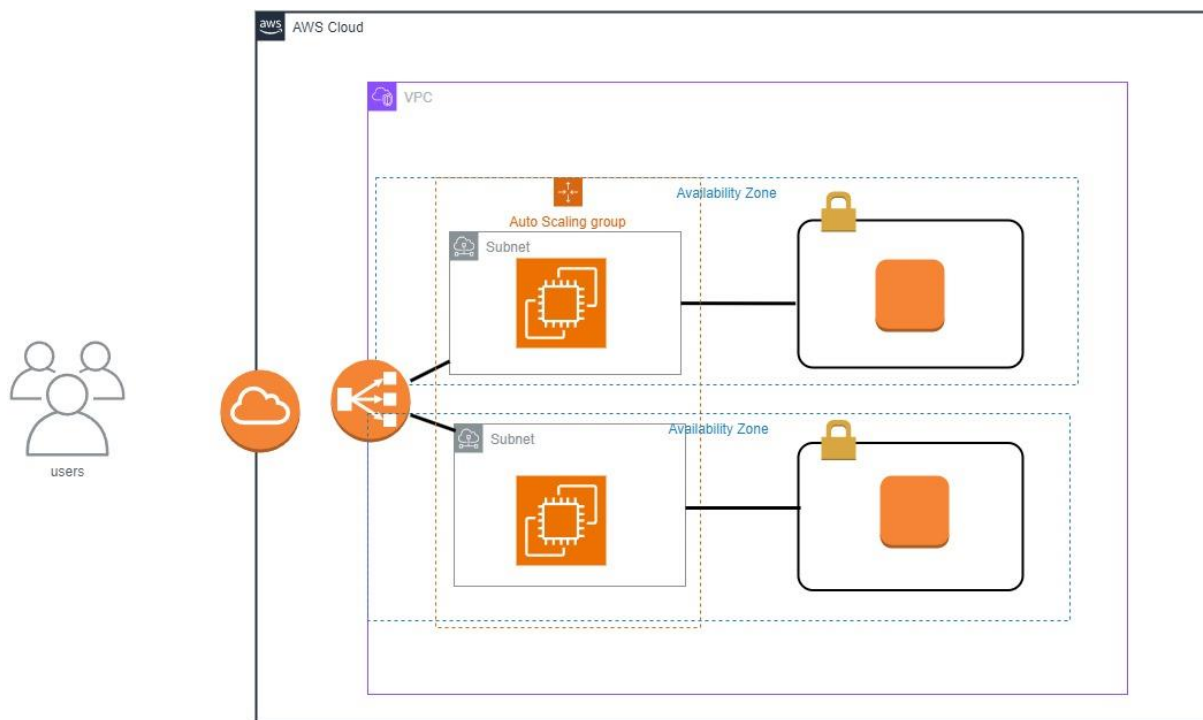
Figure 1 Installing nvm via SSH .....	5
Figure 2 Installing Node.js and running app.js .....	5
Figure 3 Initial version of app.js - Hello World. ....	6
Figure 4 Server running. ....	6
Figure 5 Updated app.js - Server instance IP address. ....	7
Figure 6 Killing the process and restarting app.js to start new version. ....	7
Figure 7 Preview of VPC configuration.....	8
Figure 8 VPC workflow complete. ....	9
Figure 9 VPC configured and active.....	10
Figure 10 VPC Route Tables.....	10
Figure 11 Security Groups .....	11
Figure 12 VPC Resource Map with a route highlighted.....	11
Figure 13 Editing Subnet Settings.....	12
Figure 14 Set IP address to Auto Assignment.....	12
Figure 15 Launch Template Success.....	13
Figure 16 Launching an Instance - part 1 – choosing an AMI.....	14
Figure 17 Launching an Instance - part 2 – Key Pair, Network Settings (Subnets), and Security Group.....	15

Figure 18 Launching an Instance - part 3 - adding a script.....	16
Figure 19 Launch Instance with key pair and Network settings.....	17
Figure 20 Before Target Group setup.....	18
Figure 21 Load balancing Target Group. ....	18
Figure 22 Load Balancer created successfully.....	19
Figure 23 Load Balancer details in browser.....	19
Figure 24 Creating an Auto Scaling group. ....	20
Figure 25 Auto Scaling - Choosing a Launch Template and Network.....	21
Figure 26 Auto Scaling -Maximum instances, SNS Notifications, and Tags.....	22
Figure 27 Auto Scaling - First Auto Scaled Instance. ....	23
Figure 28 Auto Scaling - First Instance Running. ....	23
Figure 29 Auto Scaling - Targets registered.....	24
Figure 30 Auto Scaling - Cloudwatch Metrics.....	25
Figure 31 Auto Scaling - Cloudwatch Alarms in progress. ....	25
Figure 32 Auto Scaling - Using an infinite loop to create traffic.....	26
Figure 33 Curl command generates 100 requests. ....	27
Figure 34 Adding Listeners on port 80 and port 3000. ....	27
Figure 35 Auto Scaled instances running. ....	28
Figure 36 Auto Scaled instances - Running and Terminated.....	28

## Planning: Application selection and architecture planning, including architecture diagram. AMI creation

Basic architecture:

Cloud – VPC – Availability Zones – Autoscaling group – Subnets.



I used a custom script to launch an EC2 instance, set the key pair and security group, install Node.js, upload app.js, run the app, open a web browser page to show instance information and another page to show the app in action. First 'Hello World' and later showing the Server instance IP address.

From this instance I created a custom AMI and updated the script to use this AMI for future instances which created instances with the index.html page, the monitoring script, and the application running on port 3000.

## Installation of nvm, node, and app.js

```
[ec2-user@ip-172-31-7-52 ~]$ curl -o- https://raw.githubusercontent.com/nvm-sh/nvm/v0.39.5/install.sh | bash
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed
100 15916  100 15916    0     0  423k      0 --:--:-- --:--:-- --:--:-- 431k
=> Downloading nvm as script to '/home/ec2-user/.nvm'

=> Appending nvm source string to /home/ec2-user/.bashrc
=> Appending bash_completion source string to /home/ec2-user/.bashrc
=> Close and reopen your terminal to start using nvm or run the following to use it now:

export NVM_DIR="$HOME/.nvm"
[ -s "$NVM_DIR/nvm.sh" ] && \. "$NVM_DIR/nvm.sh" # This loads nvm
[ -s "$NVM_DIR/bash_completion" ] && \. "$NVM_DIR/bash_completion" # This loads nvm bash_comp
pletion
[ec2-user@ip-172-31-7-52 ~]$ . ~/.nvm/nvm.sh
[ec2-user@ip-172-31-7-52 ~]$ nvm install 16
Downloading and installing node v16.20.2...
Downloading https://nodejs.org/dist/v16.20.2/node-v16.20.2-linux-x64.tar.xz...
##### 100.0%
Computing checksum with sha256sum
Checksums matched!
Now using node v16.20.2 (npm v8.19.4)
Creating default alias: default -> 16 (-> v16.20.2)
[ec2-user@ip-172-31-7-52 ~]$ node -v
v16.20.2
[ec2-user@ip-172-31-7-52 ~]$
```

Figure 1 Installing nvm via SSH

```
[ec2-user@ip-172-31-3-6 ~]$ node version
-bash: node: command not found
[ec2-user@ip-172-31-3-6 ~]$ nvm install node
Downloading and installing node v22.7.0...
Downloading https://nodejs.org/dist/v22.7.0/node-v22.7.0-linux-x64.tar.xz...
##### 100.0%
Computing checksum with sha256sum
Checksums matched!
Now using node v22.7.0
Creating default alias: default -> node (-> v22.7.0)
[ec2-user@ip-172-31-3-6 ~]$ node app.js
node: /lib64/libm.so.6: version `GLIBC_2.27' not found (required by node)
node: /lib64/libc.so.6: version `GLIBC_2.27' not found (required by node)
node: /lib64/libc.so.6: version `GLIBC_2.28' not found (required by node)
[ec2-user@ip-172-31-3-6 ~]$ nvm install 16
Downloading and installing node v16.20.2...
Downloading https://nodejs.org/dist/v16.20.2/node-v16.20.2-linux-x64.tar.xz...
##### 100.0%
Help ing checksum with sha256sum
Checksums matched!
Now using node v16.20.2 (npm v8.19.4)
[ec2-user@ip-172-31-3-6 ~]$ node -v
v16.20.2
[ec2-user@ip-172-31-3-6 ~]$ node app.js
Server running
```

Figure 2 Installing Node.js and running app.js






Figure 3 Initial version of app.js - Hello World.

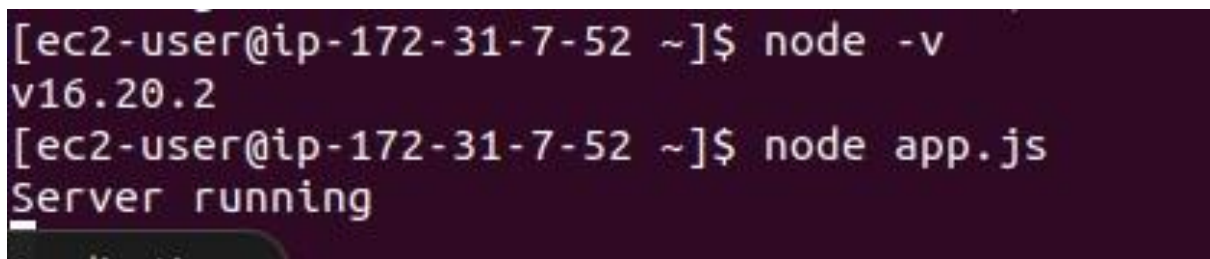


Figure 4 Server running.

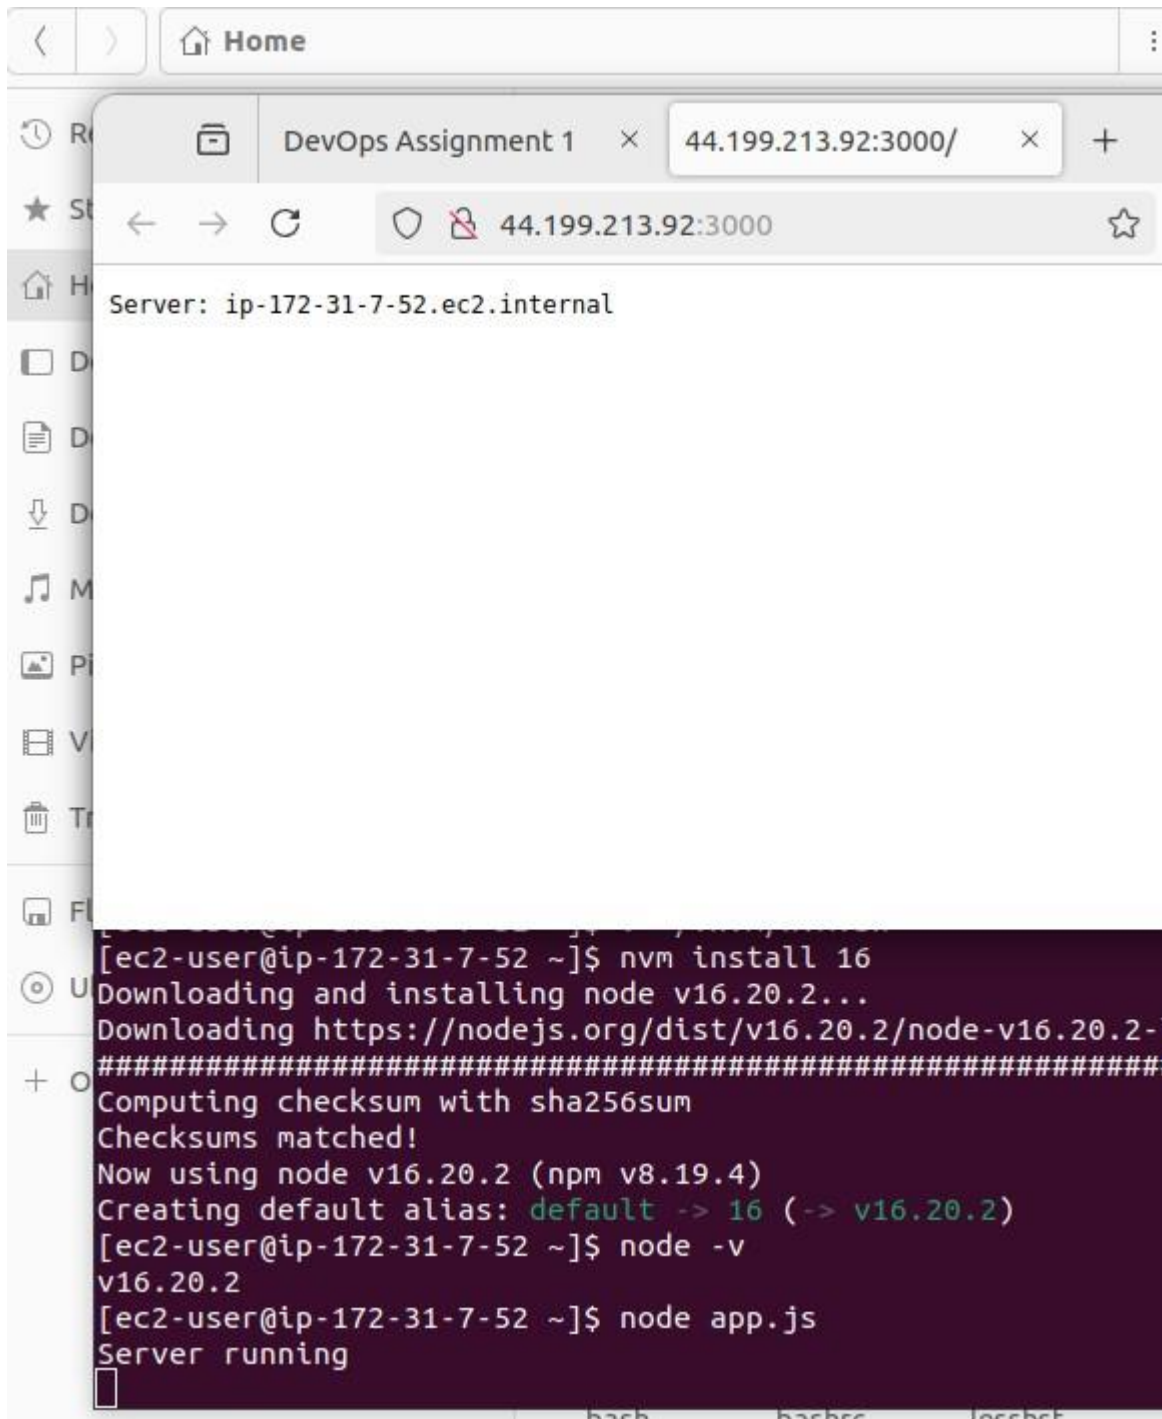


Figure 5 Updated `app.js` - Server instance IP address.

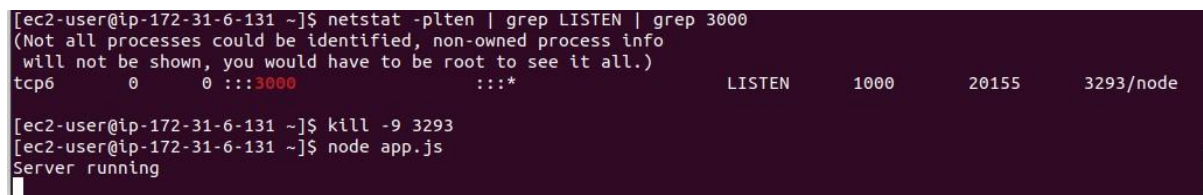


Figure 6 Killing the process and restarting `app.js` to start new version.

## VPC and load balancer implementation.

I created a VPC with three availability zones, us-east-1a, 1b, and 1c ,with public and private subnets in each.

An Elastic Load Balancer was configured to use my custom AMI and linked to an Autoscaling Group

The Load Balancer used Simple Scaling triggered by Cloudwatch alarms to increase or decrease the number of instances based on CPU usage. When usage is over 50% another instance is created, with a limit of three. When the work creating loop is terminated the low usage alarm kills unneeded instances when CPU usage drops below 30%.

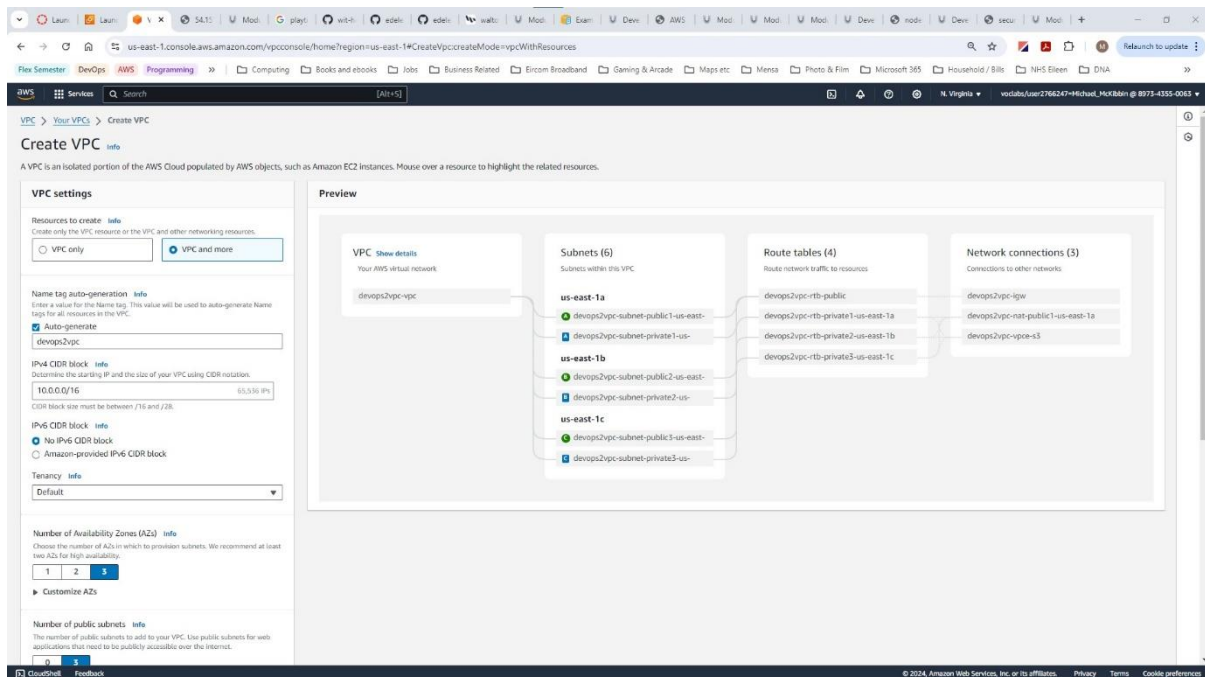


Figure 7 Preview of VPC configuration



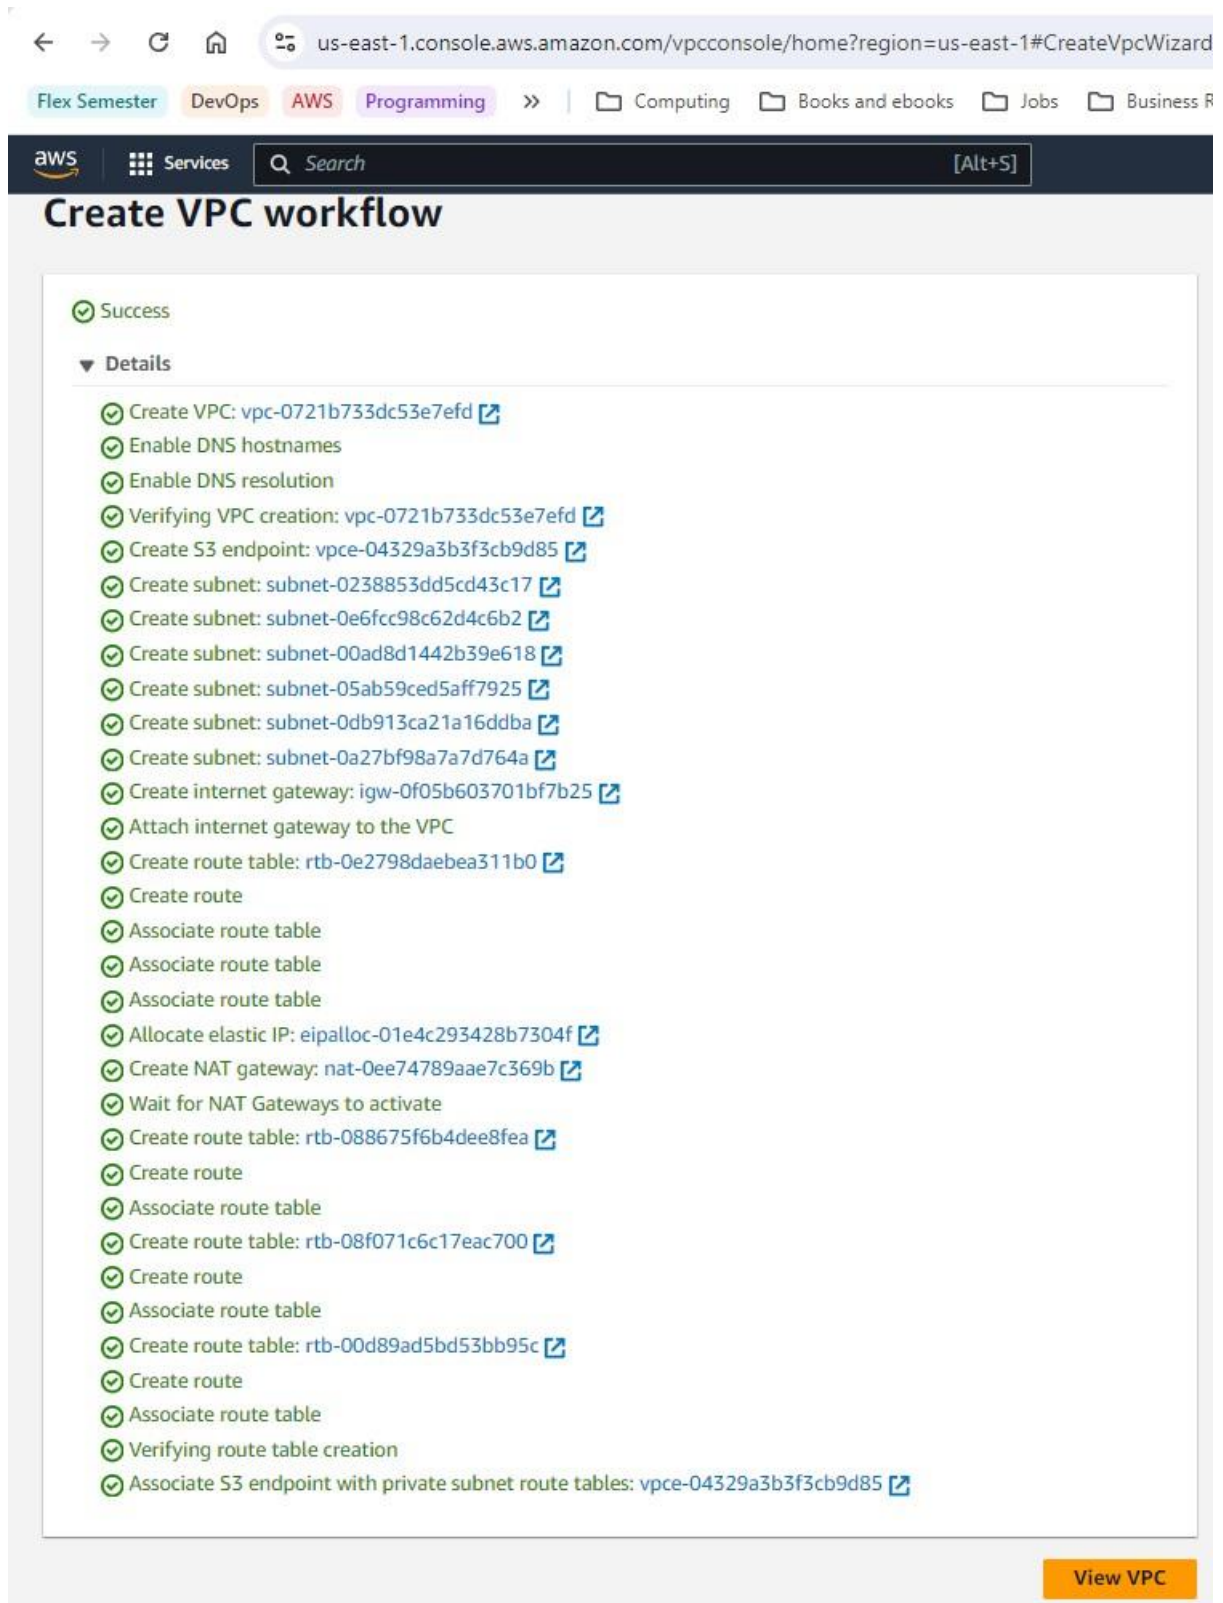


Figure 8 VPC workflow complete.

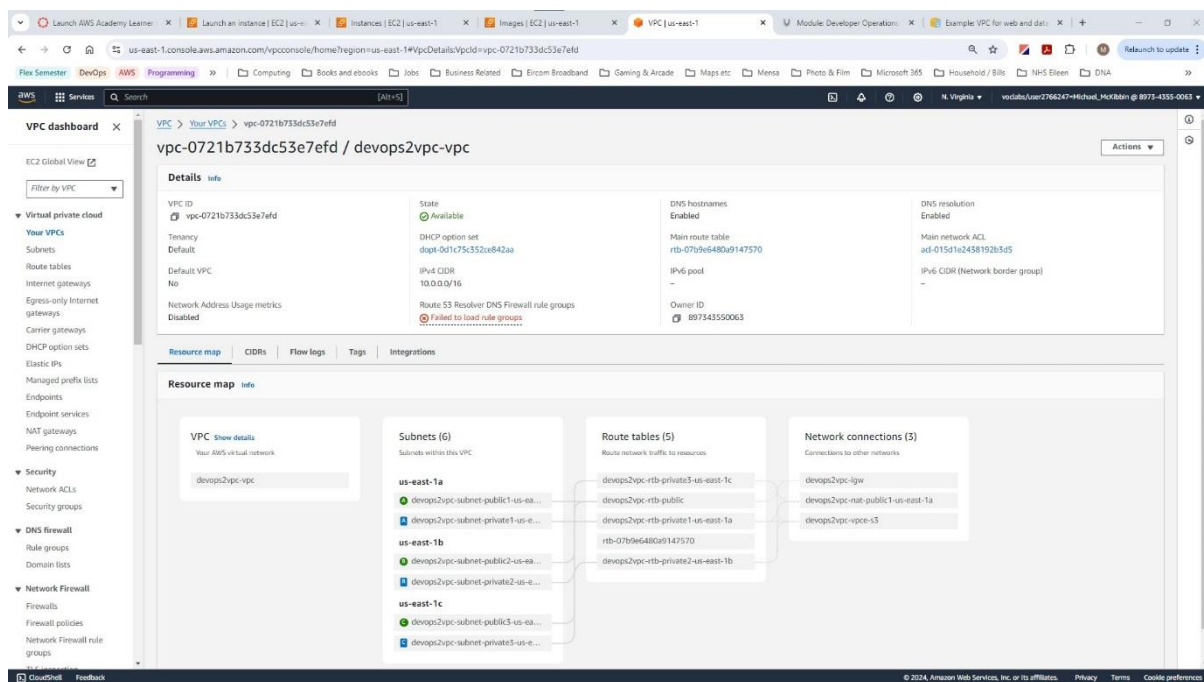


Figure 9 VPC configured and active.

VPC dashboard

Route tables (6)

Find resources by attribute or tag

Name	Route table ID	Explicit subnet associ...	Edge associations	Main	VPC	Owner ID
-	rtb-047f99367ff395853	-	-	Yes	vpc-01c538caf113ce847	897343550063
-	rtb-07b9e6480a9147570	-	-	Yes	vpc-0721b733dc53e7efd   devo...	897343550063
devops2vpc-rtb-private1-us-east-1a	rtb-088675f6b4dee8fea	subnet-05ab59ced5aff79...	-	No	vpc-0721b733dc53e7efd   devo...	897343550063
devops2vpc-rtb-private2-us-east-1b	rtb-08f071c6c17ea700	subnet-0db913ca21a16d...	-	No	vpc-0721b733dc53e7efd   devo...	897343550063
devops2vpc-rtb-private3-us-east-1c	rtb-00d89ad5bd53bb95c	subnet-0a27bf98a7a7d7...	-	No	vpc-0721b733dc53e7efd   devo...	897343550063
devops2vpc-rtb-public	rtb-0e2798daeb311b0	3 subnets	-	No	vpc-0721b733dc53e7efd   devo...	897343550063

Figure 10 VPC Route Tables.

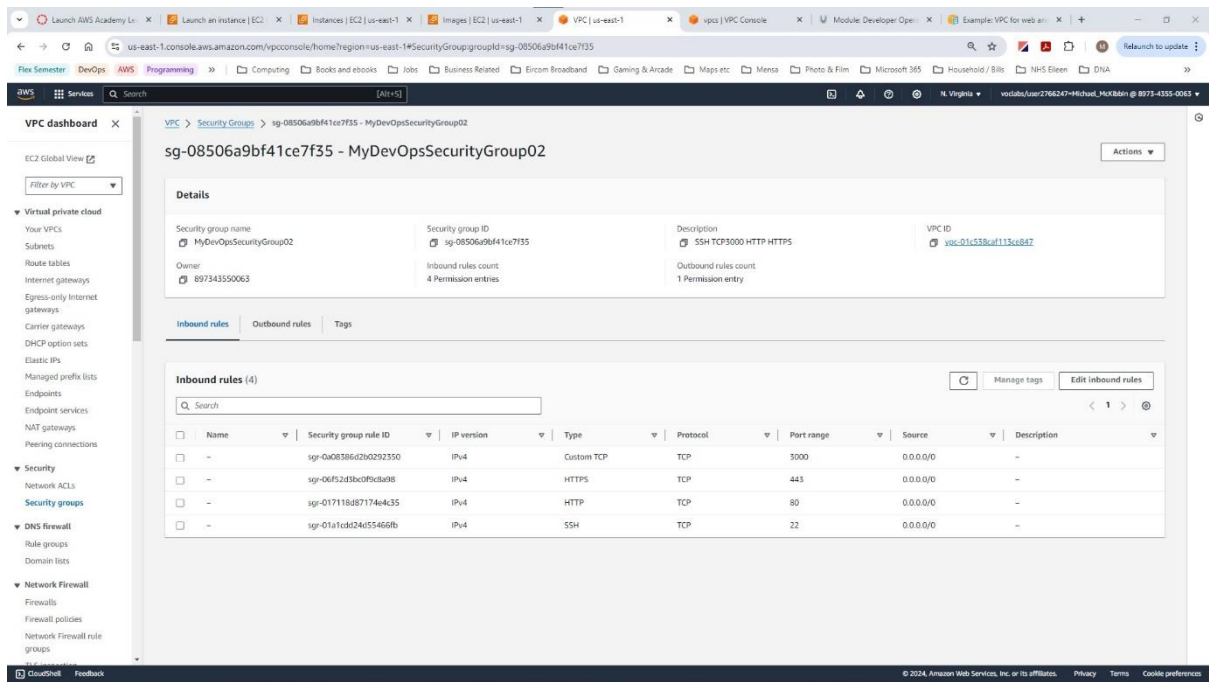


Figure 11 Security Groups

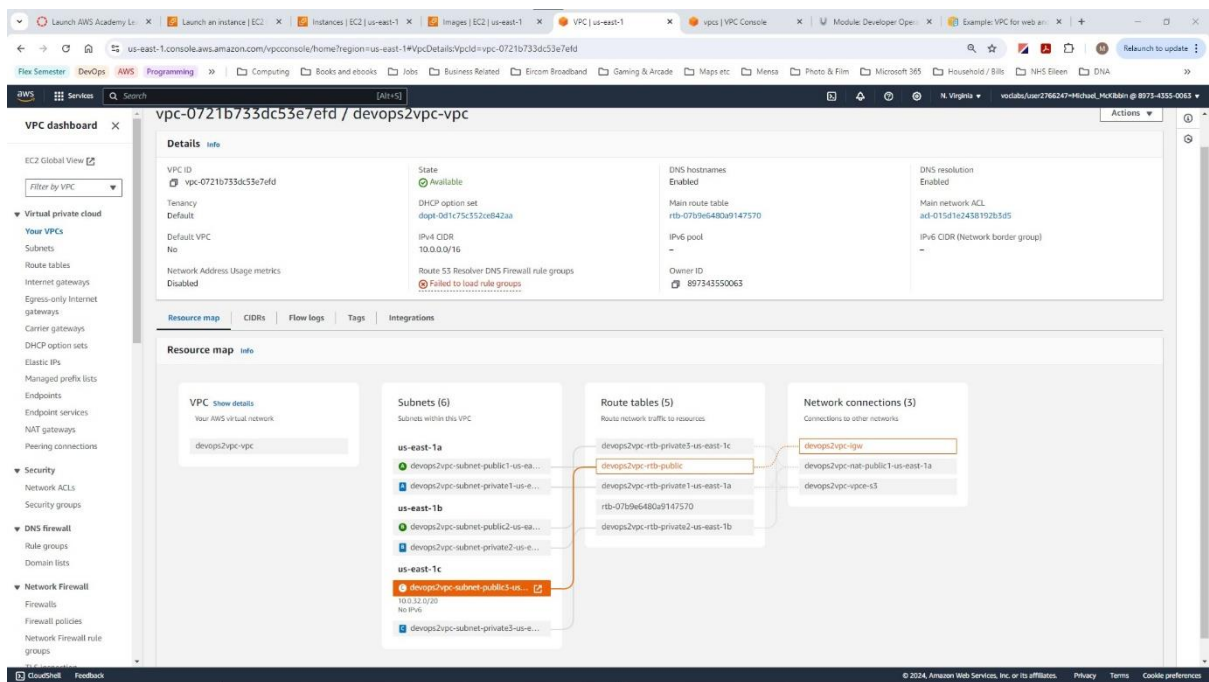


Figure 12 VPC Resource Map with a route highlighted.

VPC > Subnets > subnet-0238853dd5cd43c17 > Edit subnet settings

## Edit subnet settings [Info](#)

**Subnet**

Subnet ID	Name
subnet-0238853dd5cd43c17	devops2vpc-subnet-public1-us-east-1a

**Auto-assign IP settings** [Info](#)

Enable AWS to automatically assign a public IPv4 or IPv6 address to a new primary network interface for an instance in this subnet.

☒ Enable auto-assign public IPv4 address [Info](#)

☐ Enable auto-assign customer-owned IPv4 address [Info](#)  
Option disabled because no customer owned pools found.

**Resource-based name (RBN) settings** [Info](#)

Specify the hostname type for EC2 instances in this subnet and optional RBN DNS query settings.

☐ Enable resource name DNS A record on launch [Info](#)

☐ Enable resource name DNS AAAA record on launch [Info](#)

Hostname type [Info](#)

☐ Resource name

☒ IP name

**DNS64 settings**

Enable DNS64 to allow IPv6-only services in Amazon VPC to communicate with IPv4-only services and networks.

☐ Enable DNS64 [Info](#)

Cancel Save

Figure 13 Editing Subnet Settings.

VPC > Subnets > subnet-00ad8d1442b39e618

subnet-00ad8d1442b39e618 / devops2vpc-subnet-public3-us-east-1c Actions

Details			
Subnet ID subnet-00ad8d1442b39e618	Subnet ARN arn:aws:ec2:us-east-1:897343550063:subnet/subnet-00ad8d1442b39e618	State Available	IPv4 CIDR 10.0.32.0/20
Available IPv4 addresses 4091	IPv6 CIDR -	IPv6 CIDR association ID -	Availability Zone us-east-1c
Availability Zone ID us-east-1	Network border group us-east-1	VPC vpc-0721b733dc53e7efd   devops2vpc-vpc	Route table rtb-0e2798dae311b0   devops2vpc-rtb-public
Network ACL acl-015d1e2438192b3d5	Default subnet No	Auto-assign public IPv4 address Yes	Auto-assign IPv6 address No
Auto-assign customer-owned IPv4 address No	Customer-owned IPv4 pool -	Outpost ID -	IPv4 CIDR reservations -
IPv6 CIDR reservations -	IPv6-only No	Hostname type IP name	Resource name DNS A record Disabled
Resource name DNS AAAA record Disabled	DNS64 Disabled	Owner 897343550063	

Figure 14 Set IP address to Auto Assignment

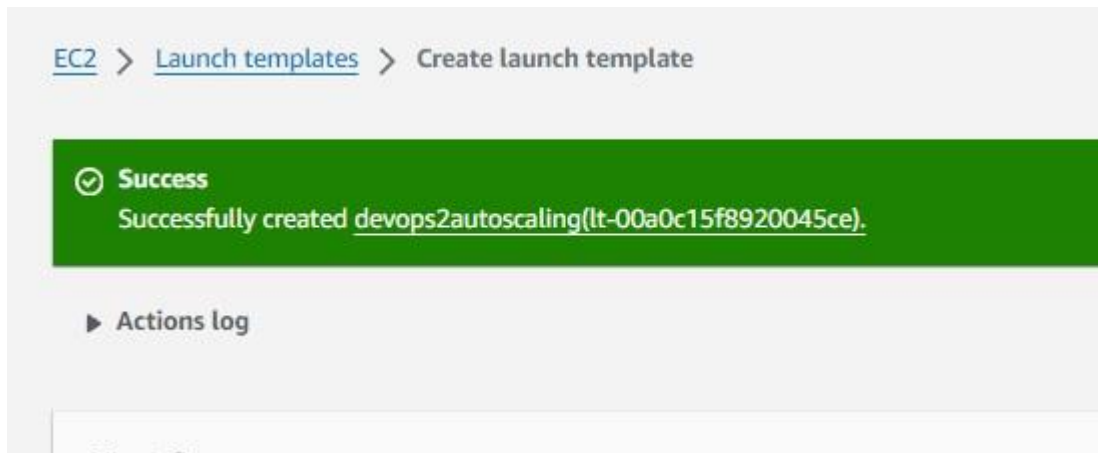


Figure 15 Launch Template Success.



[EC2](#) > [Instances](#) > Launch an instance

## Launch an instance [Info](#)

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

### Name and tags [Info](#)

Name

devops2instance 1a

[Add additional tags](#)

### ▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

AMI from catalog


Recents

**My AMIs**

Quick Start

☒ Owned by me

☐ Shared with me

  
**Browse more AMIs**  
Including AMIs from  
AWS, Marketplace and  
the Community

Amazon Machine Image (AMI)

testimage1  
ami-0ced9049444b57252  
2024-08-25T17:01:05.000Z   Virtualization: hvm   ENA enabled: true   Root device type: ebs

Description

image of working server id script

Figure 16 Launching an Instance - part 1 – choosing an AMI.

▼ Key pair (login) Info

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

key6 ▼

↻ Create new key pair

▼ Network settings Info

VPC - required Info

vpc-01c538caf113ce847 (default) ▼

172.31.0.0/16

↻

Subnet Info

subnet-040acd4f610cb94aa ▼

VPC: vpc-01c538caf113ce847 Owner: 897343550063 Availability Zone: us-east-1a  
Zone type: Availability Zone IP addresses available: 4090 CIDR: 172.31.16.0/20

↻ Create new subnet

Auto-assign public IP Info

Enable ▼

Additional charges apply when outside of free tier allowance

Firewall (security groups) Info

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☐ Create security group
 ☒ Select existing security group

Common security groups Info

Select security groups ▼

MyDevOpsSecurityGroup02 sg-08506a9bf41ce7f35 X

VPC: vpc-01c538caf113ce847

↻ Compare security group rules

Security groups that you add or remove here will be added to or removed from all your network interfaces.

▶ Advanced network configuration

Figure 17 Launching an Instance - part 2 – Key Pair, Network Settings (Subnets), and Security Group.

The screenshot shows the AWS Management Console interface for launching an EC2 instance. The main configuration area on the left includes several dropdown menus for metadata settings: 'Metadata accessible', 'Metadata IPv6 endpoint', 'Metadata version', 'Metadata response hop limit', and 'Allow tags in metadata'. A yellow warning box states: 'EC2 recommends using metadata version 2 unless you explicitly require metadata version 1.' Below these, there is a 'User data - optional' section with a 'Choose file' button and a text area containing the script: `#!/bin/bash` and `su - ec2-user -c 'node app.js'`. A checkbox at the bottom indicates 'User data has already been base64 encoded'. The right pane shows a 'Summary' section with fields for 'Number of instances' (set to 1), 'Software Image (AMI)' (ami-0ced9049444b57252), 'Virtual server type (instance type)' (t2.nano), 'Firewall (security group)' (MyDevOpsSecurityGroup02), and 'Storage (volumes)' (1 volume(s) - 8 GiB). A blue information box highlights the 'Free tier' benefits. At the bottom right, there are 'Cancel', 'Launch instance', and 'Review commands' buttons.

Figure 18 Launching an Instance - part 3 - adding a script.

▼ Instance type [Info](#) | [Get advice](#)

Instance type

t2.nano

Family: t2 1 vCPU 0.5 GiB Memory Current generation: true  
On-Demand Linux base pricing: 0.0058 USD per Hour  
On-Demand SUSE base pricing: 0.0058 USD per Hour  
On-Demand Windows base pricing: 0.0081 USD per Hour

☐ All generations

[Compare instance types](#)

[Additional costs apply for AMIs with pre-installed software](#)

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - *required*

key6

[Create new key pair](#)

▼ Network settings [Info](#)

VPC - *required* [Info](#)

vpc-01c538caf113ce847

172.31.0.0/16

(default) ▼

[Create new VPC](#)

Subnet [Info](#)

subnet-07edea8fec1c2bf64

VPC: vpc-01c538caf113ce847 Owner: 897343550063 Availability Zone: us-east-1b  
Zone type: Availability Zone IP addresses available: 4089 CIDR: 172.31.32.0/20

[Create new subnet](#)

Auto-assign public IP [Info](#)

Enable

[Additional charges apply when outside of free tier allowance](#)

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

Figure 19 Launch Instance with key pair and Network settings.

## Auto-scaling implementation and demonstration of scaling activity based on CloudWatch alarm.

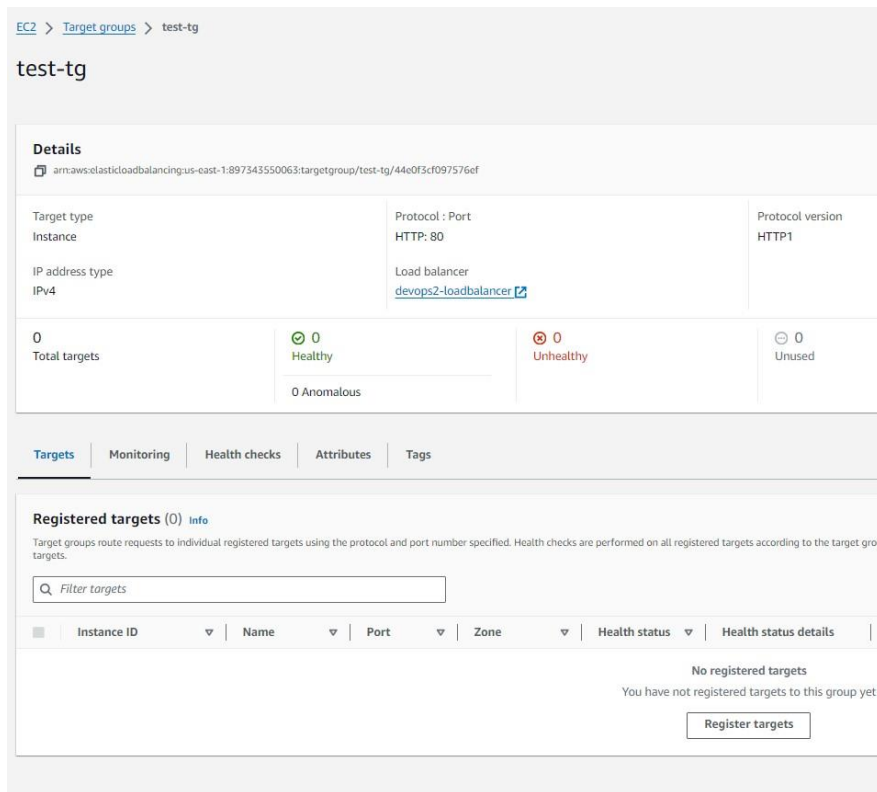


Figure 20 Before Target Group setup.

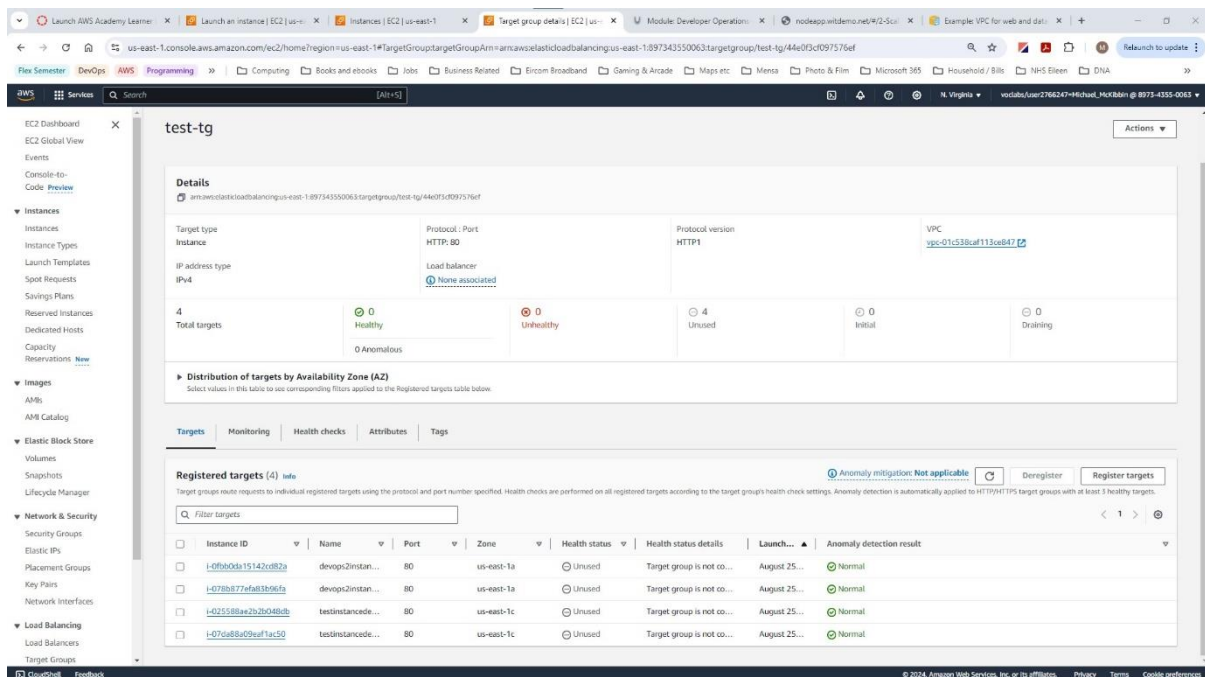


Figure 21 Load balancing Target Group.



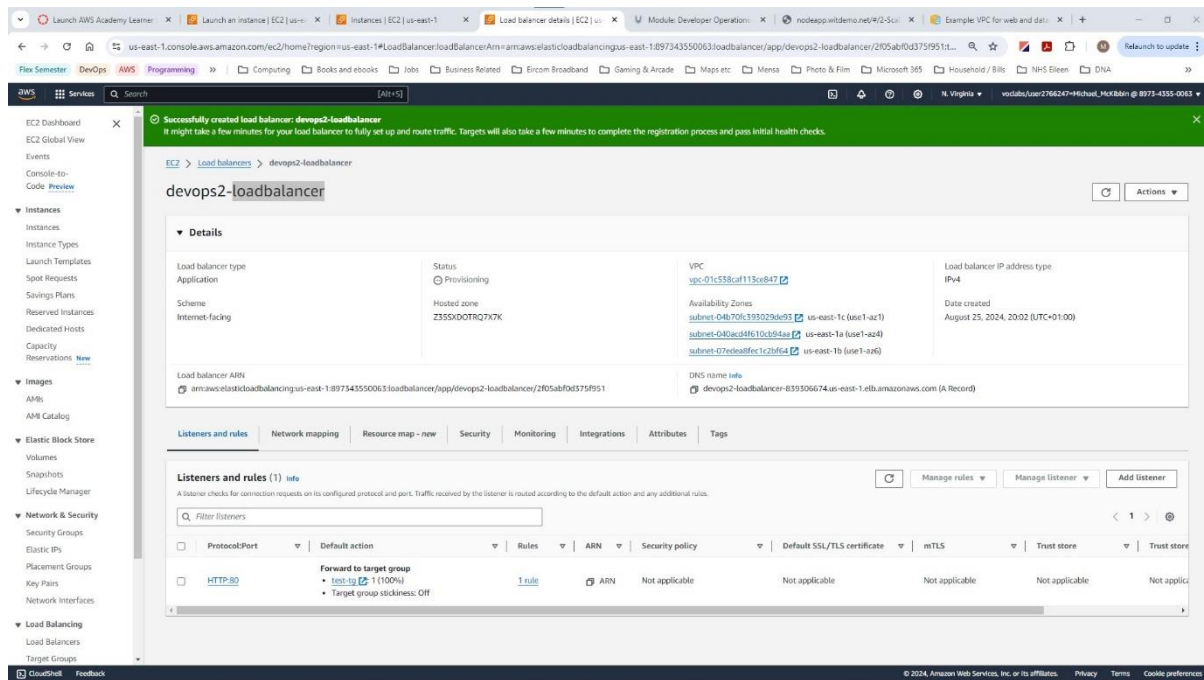



Figure 22 Load Balancer created successfully.



## Assignment 2: DevOps\_2

AMI: ami-07761f3ae34c4478d

Instance ID: i-016bf77858075e1f1

Instance Type: t2.nano

Instance IP: 54.157.122.46

Security Group: MyDevOpsSecurityGroup02

Availability Zone: use1-az1

SETU logo

Figure 23 Load Balancer details in browser.

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1  
[Choose launch template](#)

Step 2  
[Choose instance launch options](#)

Step 3 - optional  
**Configure advanced options**

Step 4 - optional  
[Configure group size and scaling](#)

Step 5 - optional  
[Add notifications](#)

Step 6 - optional  
[Add tags](#)

Step 7  
[Review](#)

## Configure advanced options - optional Info

Integrate your Auto Scaling group with other services to distribute network traffic across multiple servers using a load balancer or to establish service-to-service communications using VPC Lattice. You can also set options that give you more control over health check replacements and monitoring.

### Load balancing Info

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

☐ No load balancer  
Traffic to your Auto Scaling group will not be fronted by a load balancer.

☒ Attach to an existing load balancer  
Choose from your existing load balancers.

☐ Attach to a new load balancer  
Quickly create a basic load balancer to attach to your Auto Scaling group.

### Attach to an existing load balancer

Select the load balancers that you want to attach to your Auto Scaling group.

☒ Choose from your load balancer target groups  
This option allows you to attach Application, Network, or Gateway Load Balancers.

☐ Choose from Classic Load Balancers

Existing load balancer target groups  
Only instance target groups that belong to the same VPC as your Auto Scaling group are available for selection.

Select target groups

test-tg | HTTP  
Application Load Balancer: devops2-loadbalancer

×

### VPC Lattice integration options Info

Figure 24 Creating an Auto Scaling group.

[Linkedin: michaelkevinmckibbin](#)

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1  
[Choose launch template](#)

Step 2  
[Choose instance launch options](#)

Step 3 - optional  
[Configure advanced options](#)

Step 4 - optional  
[Configure group size and scaling](#)

Step 5 - optional  
[Add notifications](#)

Step 6 - optional  
[Add tags](#)

Step 7  
**Review**

## Review Info

### Step 1: Choose launch template Edit

**Group details**

Auto Scaling group name  
devops2autoscale

**Launch template**

Launch template	Version	Description
autoscalelaunchtemplate1 <a href="#">↗</a> lt-0979327f778f62c91	Default	devops2

### Step 2: Choose instance launch options Edit

**Network**

Network  
VPC  
vpc-01c538caf113ce847 [↗](#)

Availability Zone	Subnet	
us-east-1a	subnet-040acd4f610cb94aa <a href="#">↗</a>	172.31.16.0/20
us-east-1b	subnet-07edea8fec1c2bf64 <a href="#">↗</a>	172.31.32.0/20
us-east-1c	subnet-04b70fc393029de93 <a href="#">↗</a>	172.31.0.0/20

**Instance type requirements**

This Auto Scaling group will adhere to the launch template.

### Step 3: Configure advanced options Edit

**Load balancing**

Load balancer 1

Figure 25 Auto Scaling - Choosing a Launch Template and Network.

Desired capacity 1	Desired capacity type Units (number of instances)
-----------------------	--

### Scaling

Minimum desired capacity 1	Maximum desired capacity 3
Target tracking policy -	

### Instance maintenance policy

Replacement behavior No policy	Min healthy percentage -	Max healthy percentage -
-----------------------------------	-----------------------------	-----------------------------

### Instance scale-in protection

Instance scale-in protection

☐ Enable instance protection from scale in

Step 5: Add notifications Edit

### Notifications

Notification 1 SNS Topic devops2test (mmckibbin1@gmail.com)	Event types <input checked="" type="checkbox"/> Launch <input checked="" type="checkbox"/> Terminate <input checked="" type="checkbox"/> Fail to launch <input checked="" type="checkbox"/> Fail to terminate
---	---

Step 6: Add tags Edit

### Tags (1)

Key	Value	Tag new instances
Name	devops2 auto scaled instance	Yes

Cancel
Previous
Create Auto Scaling group

Figure 26 Auto Scaling -Maximum instances, SNS Notifications, and Tags.

Instances (6) [Info](#) Last updated less than a minute ago

Find Instance by attribute or tag (case-sensitive) All states ▼

Instance state = running X Clear filters

<input type="checkbox"/>	Name <a href="#">↗</a>	Instance ID	Instance state <a href="#">▼</a>	Instance type <a href="#">▼</a>	Status check
<input type="checkbox"/>	devops2 auto scaled instance	i-0df3bd6fc454514c4	Running <a href="#">🔍</a> <a href="#">🔍</a>	t2.nano	⌚ Initializing
<input type="checkbox"/>	devops2withscript	i-01f64a566f3404da8	Running <a href="#">🔍</a> <a href="#">🔍</a>	t2.nano	✔ 2/2 checks passed..
<input type="checkbox"/>	Devops2Hello	i-0f27c3281cae8d7a7	Running <a href="#">🔍</a> <a href="#">🔍</a>	t2.nano	✔ 2/2 checks passed..
<input type="checkbox"/>	MyInstance1	i-0ead3a853d8a1e5e8	Running <a href="#">🔍</a> <a href="#">🔍</a>	t2.nano	✔ 2/2 checks passed..
<input type="checkbox"/>	Week4 Instance 1	i-0b0707a632ede78fe	Running <a href="#">🔍</a> <a href="#">🔍</a>	t2.nano	✔ 2/2 checks passed..
<input type="checkbox"/>	MMCKDevOps2	i-0b09d47b249820ee3	Running <a href="#">🔍</a> <a href="#">🔍</a>	t2.nano	✔ 2/2 checks passed..

Figure 27 Auto Scaling - First Auto Scaled Instance.

← → ↻ 🏠 ⚠ Not secure devops2-loadbalancer-839306674.us-east-1.elb.amazonaws.com

Flex Semester DevOps AWS Programming >> | 📁 Computing 📁 Books and ebooks 📁 Jobs

## Assignment 2: DevOps\_2

AMI: ami-07761f3ae34c4478d

Instance ID: i-016bf77858075e1f1

Instance Type: t2.nano

Instance IP: 54.157.122.46


Security Group: MyDevOpsSecurityGroup02


Availability Zone: use1-az1

Figure 28 Auto Scaling - First Instance Running.





## test-tg


**Details**  
 `arn:aws:elasticloadbalancing:us-east-1:897343550063:targetgroup/test-tg/44e0f3cf097576ef`

Target type Instance	Protocol : Port HTTP: 80	Protocol HTTP
IP address type IPv4	Load balancer <a href="#">devops2-loadbalancer</a> 	

1  
Total targets

 1  
Healthy  
0 Anomalous

 0  
Unhealthy

 Un

► **Distribution of targets by Availability Zone (AZ)**  
 Select values in this table to see corresponding filters applied to the Registered targets table below.

**Targets** | Monitoring | Health checks | Attributes | Tags

**Registered targets (1)** [Info](#)  
 Target groups route requests to individual registered targets using the protocol and port number specified. Health checks are performed on all registered targets with at least 3 healthy targets.


<input type="checkbox"/>	Instance ID ▾	Name ▾	Port ▾	Zone ▾	Health status ▾
<input type="checkbox"/>	<a href="#">i-0df3bd6fc454514c4</a>	devops2 auto s...	80	us-east-1a	 Healthy

Figure 29 Auto Scaling - Targets registered.

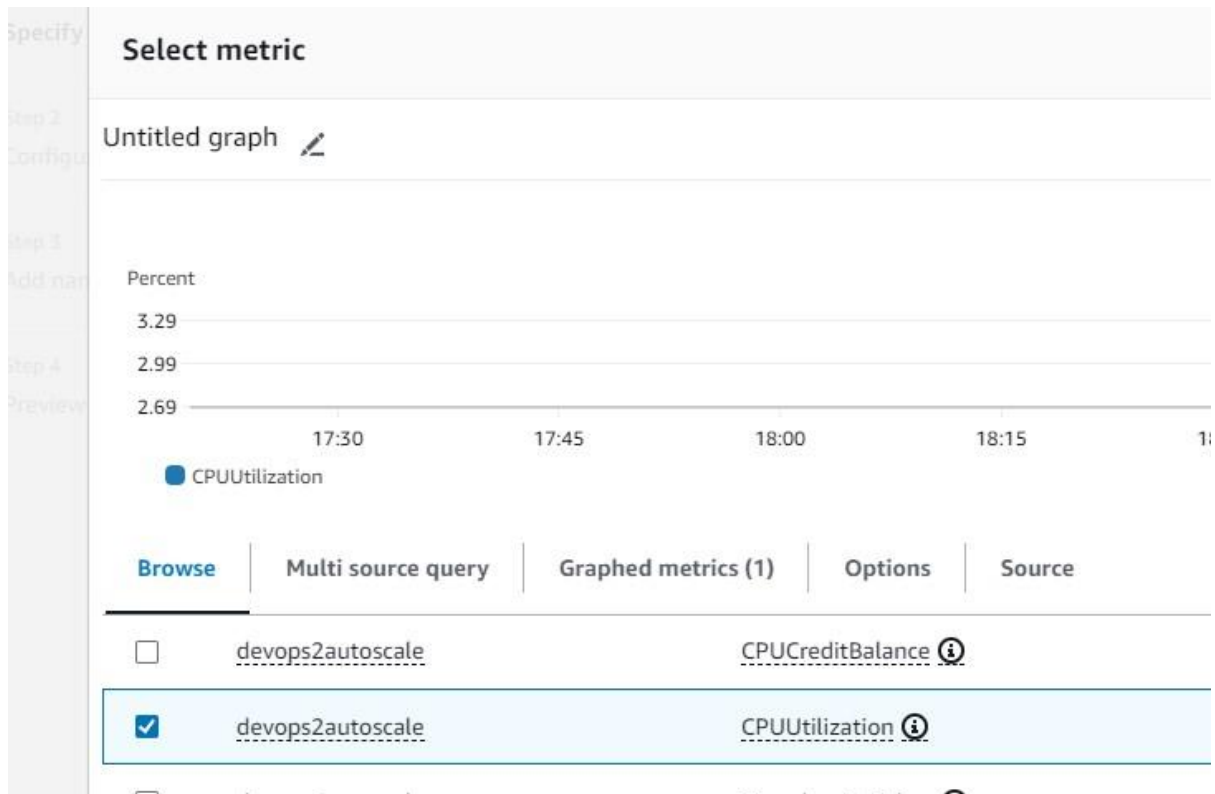


Figure 30 Auto Scaling - Cloudwatch Metrics.

CloudWatch

CloudWatch > Alarms

Favorites and recents

Dashboards

Alarms

In alarm

All alarms

Billing

Logs

Log groups

Alarms (3)

☐ Hide Auto Scaling alarms

Clear selection Create composite alarm Actions

Search Alarm state: Any Alarm type: Any Actions status: Any

<input type="checkbox"/>	Name	State	Last state update (UTC)	Conditions	Actions
<input type="checkbox"/>	TargetTracking-devops2autoscale-AlarmLow-8389be51-3763-4f42-9709-7d4c6fc4f83e	In alarm	2024-08-25 20:32:33	CPUUtilization < 35 for 15 datapoints within 15 minutes	Actions enabled
<input type="checkbox"/>	TargetTracking-devops2autoscale-AlarmHigh-c58812e2-0851-4150-9fd4-0085374fc39a	OK	2024-08-25 20:32:10	CPUUtilization > 50 for 3 datapoints within 3 minutes	Actions enabled
<input type="checkbox"/>	devops2 cloudwatch high cpu	OK	2024-08-25 20:30:20	CPUUtilization > 40 for 1 datapoints within 1 minute	Actions enabled

Figure 31 Auto Scaling - Cloudwatch Alarms in progress.



```

root@ip-172-31-18-111:/var/log/httpd
GNU nano 2.9.8 access_log
86.43.222.44 - - [25/Aug/2024:16:53:55 +0000] "GET / HTTP/1.1" 200 780 "-" "Mozilla/5.0 (X11; Ubuntu; L$
86.43.222.44 - - [25/Aug/2024:16:53:56 +0000] "GET /favicon.ico HTTP/1.1" 404 196 "http://54.157.122.46$
86.43.222.44 - - [25/Aug/2024:16:57:11 +0000] "GET / HTTP/1.1" 200 780 "-" "Mozilla/5.0 (Windows NT 10.$
86.43.222.44 - - [25/Aug/2024:16:57:11 +0000] "GET /favicon.ico HTTP/1.1" 404 196 "http://54.157.122.46$
172.31.0.176 - - [25/Aug/2024:20:04:57 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:04:57 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.0.176 - - [25/Aug/2024:20:05:20 +0000] "GET /ReportServer HTTP/1.1" 404 196 "-" "Mozilla/5.0 zgr$
172.31.0.176 - - [25/Aug/2024:20:05:27 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:05:27 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.0.176 - - [25/Aug/2024:20:05:57 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:05:57 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.0.176 - - [25/Aug/2024:20:06:09 +0000] "GET /.git/config HTTP/1.1" 404 196 "-" "Mozilla/5.0 (MacS$
172.31.0.176 - - [25/Aug/2024:20:06:18 +0000] "GET / HTTP/1.1" 304 - "-" "Mozilla/5.0 (Windows NT 10.0;$
172.31.0.176 - - [25/Aug/2024:20:06:27 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:06:27 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.0.176 - - [25/Aug/2024:20:06:57 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:06:57 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.0.176 - - [25/Aug/2024:20:07:27 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:07:27 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:07:55 +0000] "GET / HTTP/1.1" 200 780 "-" "Mozilla/5.0 (Macintosh; Int$
172.31.0.176 - - [25/Aug/2024:20:07:57 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:07:57 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:08:13 +0000] "GET /favicon.ico HTTP/1.1" 404 196 "-" "Mozilla/5.0 (MacS$
172.31.18.74 - - [25/Aug/2024:20:08:14 +0000] "GET /robots.txt HTTP/1.1" 404 196 "-" "Mozilla/5.0 (Maci$
172.31.18.74 - - [25/Aug/2024:20:08:14 +0000] "GET /sitemap.xml HTTP/1.1" 404 196 "-" "Mozilla/5.0 (MacS$
172.31.18.74 - - [25/Aug/2024:20:08:14 +0000] "GET /config.json HTTP/1.1" 404 196 "-" "Mozilla/5.0 (MacS$
172.31.0.176 - - [25/Aug/2024:20:08:27 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.18.74 - - [25/Aug/2024:20:08:27 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"
172.31.0.176 - - [25/Aug/2024:20:08:57 +0000] "GET / HTTP/1.1" 200 780 "-" "ELB-HealthChecker/2.0"

```

Figure 33 Curl command generates 100 requests.

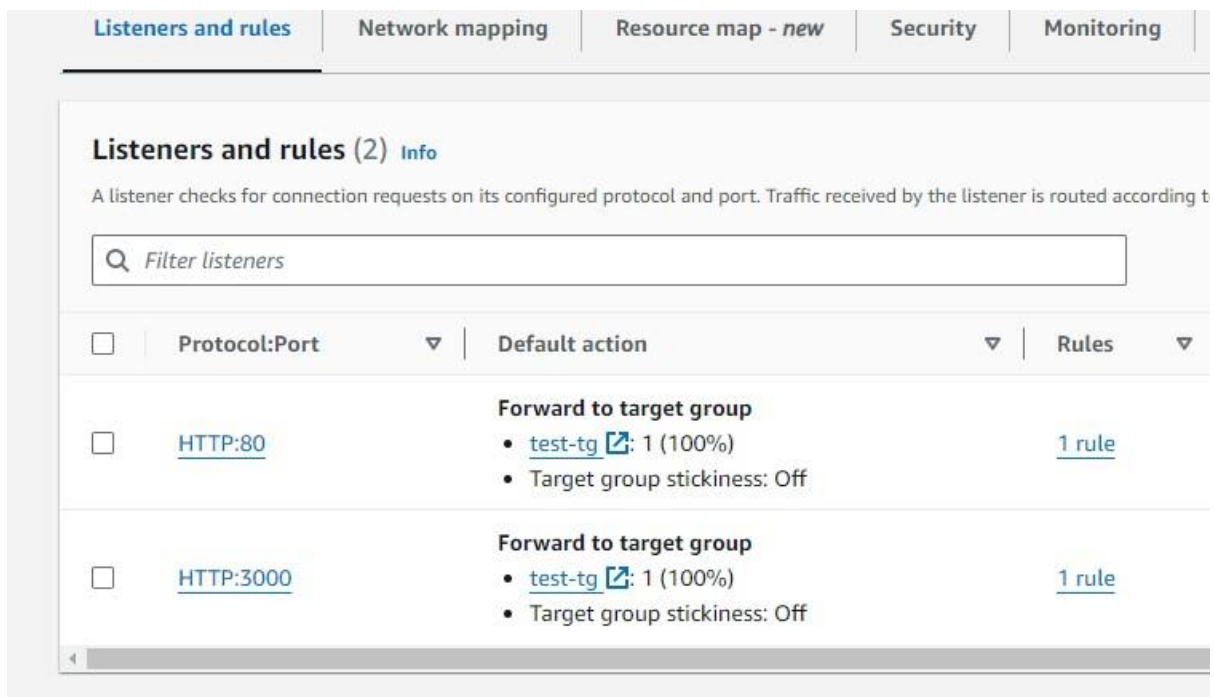


Figure 34 Adding Listeners on port 80 and port 3000.

<input type="checkbox"/>	devops2Autoscale	i-00928be326140d4c9	✓ Running ⓘ ⓘ	t2.nano	✓ 2/2 checks passed..
<input type="checkbox"/>	devops2Autoscale	i-04929e93f7d1b95eb	✓ Running ⓘ ⓘ	t2.nano	✓ 2/2 checks passed..
<input type="checkbox"/>	devops2Autoscale	i-0df43dcdbd33e3457a	✓ Running ⓘ ⓘ	t2.nano	⌚ Initializing

Figure 35 Auto Scaled instances running.

<input checked="" type="checkbox"/>	devops2Autoscale	i-00928be326140d4c9	✓ Running ⓘ ⓘ	t2.nano
<input type="checkbox"/>	devops2Autoscale	i-04929e93f7d1b95eb	✓ Running ⓘ ⓘ	t2.nano
<input type="checkbox"/>	devops2Autoscale	i-0a2c3e8ce60a15b7c	⊖ Terminated ⓘ ⓘ	t2.nano
<input type="checkbox"/>	devops2Autoscale	i-017644d813d549ac7	⊖ Terminated ⓘ ⓘ	t2.nano
<input type="checkbox"/>	devops2Autoscale	i-0b233f62b39c5000a	⊖ Terminated ⓘ ⓘ	t2.nano
<input type="checkbox"/>	devops2Autoscale	i-00b95b1034097f04d	⊖ Terminated ⓘ ⓘ	t2.nano
<input type="checkbox"/>	devops2Autoscale	i-03f2fd789f621c418	⊖ Terminated ⓘ ⓘ	t2.nano
<input type="checkbox"/>	devops2Autoscale	i-0565144ba0a2c68fa	⊖ Terminated ⓘ ⓘ	t2.nano
<input type="checkbox"/>	devops2Autoscale	i-0df43dcdbd33e3457a	✓ Running ⓘ ⓘ	t2.nano

Figure 36 Auto Scaled instances - Running and Terminated.



## Monitoring: Own script.

I used my own python program, `devops_2.py`, adapted from assignment 1, to launch an instance and set up Cloudwatch alarms. I also used it to load and run further scripts, `monitor.sh`, and `memv2.sh`, to monitor and report on the custom metrics on the servers and push these to CloudWatch.

## APPENDIX A: Automated instance setup.

Including launching and configuring instance based on AMI.

Setting Key Pair and Security group.

Create webpage with content and display in a browser window.

Get instance details.

Install and run monitoring scripts – see Appendices B, and C.

Install Node.js, copy app.js, and run app in a browser window.

Set up Cloudwatch Alarms

```
import boto3
import time
import webbrowser
import random
import string
import urllib
import os
from cgitb import html
from ipaddress import ip_address
from os import system
from urllib import response
from botocore.exceptions import ClientError
import json

# Launch New EC2 Instance
# Configure instance settings
# Set up EC2 website
# Download and install Apache web server
# Create index.html and add content
# Get meta data
# Get image
# copy index.html to local drive

ec2 = boto3.resource("ec2")
try:
    print("\nCreating new EC2 instance\n")

    new_ec2_instance = ec2.create_instances(
        # Amazon Linux2 AMI (check current available AMIs before assignment
        submission)
        ImageId="ami-0ced9049444b57252",
```

```

# How many instances to launch. Min and Max
MinCount=1,
MaxCount=3,
# Instance type (t2.nano)
InstanceType="t2.nano",
# Tag the instance
TagSpecifications=[
    {
        "ResourceType": "instance",
        "Tags": [
            {"Key": "Name", "Value": "devops2Autoscale"},
            {"Key": "Owner", "Value": "mmckibbin"},
        ],
    },
],
# MyDevOpsSecurityGroup01 Security group ID
SecurityGroupIds=["sg-08506a9bf41ce7f35"],
KeyName="key6",
UserData="""#!/bin/bash
    yum install httpd -y
    systemctl enable httpd
    systemctl start httpd

    echo "Content-type: text/html"
    echo '<html>' > index.html
    echo '<head>' >> index.html
    echo '<title>DevOps Assignment 1</title>' >> index.html
    echo '</head>' >> index.html
    echo '<body>' >> index.html

    echo '<p style="font-family:Helvetica, sans-serif; font-size:200%;">Assignment 2: DevOps_2</p>' >> index.html

    echo '<p style="font-family:Helvetica, sans-serif; font-size:150%;">AMI: ' >> index.html
    echo $(curl http://169.254.169.254/latest/meta-data/ami-id) >> index.html
    echo '</p>' >> index.html

    echo '<p style="font-family:Helvetica, sans-serif; font-size:150%;">Instance ID: ' >> index.html
    echo $(curl http://169.254.169.254/latest/meta-data/instance-id) >> index.html
    echo '</p>' >> index.html

    echo '<p style="font-family:Helvetica, sans-serif; font-size:150%;">Instance Type: ' >> index.html

```

```

        echo $(curl http://169.254.169.254/latest/meta-data/instance-type)
>> index.html
        echo '</p>' >> index.html

        echo '<p style="font-family:Helvetica, sans-serif; font-
size:150%;">Instance IP: ' >> index.html
        echo $(curl http://169.254.169.254/latest/meta-data/public-ipv4)
>> index.html
        echo '</p>' >> index.html

        echo '<p style="font-family:Helvetica, sans-serif;">Security
Group: ' >> index.html
        echo $(curl http://169.254.169.254/latest/meta-data/security-
groups) >> index.html
        echo '</p>' >> index.html

        echo '<p style="font-family:Helvetica, sans-serif;">Availability
Zone: ' >> index.html
        echo $(curl http://169.254.169.254/latest/meta-
data/placement/availability-zone-id ) >> index.html
        echo '</p>' >> index.html

        echo ' logo.jpg ' >>
index.html
        echo '</body>' >> index.html
        echo '</html>' >> index.html

        cp index.html /var/www/html/index.html
    """
)

except Exception as e:
    print("Error! \nThe EC2 creation process has encountered an error.\n")
    print(e)
    errorfile = open("error.log", "w")
    errorfile.write(str(e))
    errorfile.close()
    print("See error.log for details.")

else:
    # Print instance ID, type, & state
    print(
        "\nNew EC2 instance created successfully!"
        + "\n[ID: "
        + new_ec2_instance[0].id
        + "]"
        + "\n[Type: "
        + new_ec2_instance[0].instance_type

```

```

    + "]"
    +
    #'\\n[Region: ' + new_ec2_instance[0].region['Name'] + ']' +
    "\\n[Current state: "
    + new_ec2_instance[0].state["Name"]
    + "]"
)

# Check instance state every 5 seconds.
print("\\nWaiting for instance to run...")
while new_ec2_instance[0].state["Name"] != "running":
    time.sleep(5)
    new_ec2_instance[0].reload()

# Print instance state & public ip address
print(
    "\\n\\n[Current state: "
    + new_ec2_instance[0].state["Name"]
    + "]\\n"
    + "[Public IP: "
    + new_ec2_instance[0].public_ip_address
    + "]\\n"
)

# Wait x seconds for webserver to initialise
print("\\n\\nAllowing time for web server to initialise...")
time.sleep(60)
print("\\nOpening webpage at: " + new_ec2_instance[0].public_ip_address)
print("\\n\\n")
ip_address = new_ec2_instance[
    0
].public_ip_address # Set variable to instance public IP
webbrowser.open(
    "http://" + new_ec2_instance[0].public_ip_address
) # Open web browser to instance public IP

# 6. Monitoring
# Out of sequence numerically as connection timed out when running later
on, after s3 setup (Sections 4 & 5)
try:
    time.sleep(30) # wait a bit...
    # set keypair permissions for ssh access
    print("\\nSet keypair permission")
    system("chmod 400 key6.pem")
    print("\\nDone.")
    print("\\n")

# copy monitoring scripts to instance, run them.

```

```

print("\nCopying monitor.sh to ec2 instance")
system(
    f"scp -o StrictHostKeyChecking=no -i key6.pem monitor.sh ec2-
user@{ip_address}:".
)
print("\nDone.")
print("\n")
print("\nSet permissions on monitor.sh")
system(f"ssh -i key6.pem ec2-user@{ip_address} 'chmod 700 monitor.sh'")
print("\nDone.")
print("\n")
print("\nRun monitor.sh (on ec2 instance)")
system(f"ssh -i key6.pem ec2-user@{ip_address} './monitor.sh'")
print("\nend of monitoring script")
print("\n")

print("\nCopying memv2.sh to ec2 instance")
system(
    f"scp -o StrictHostKeyChecking=no -i key6.pem memv2.sh ec2-
user@{ip_address}:".
)
print("\nDone.")
print("\n")
print("\nSet permissions on memv2.sh")
system(f"ssh -i key6.pem ec2-user@{ip_address} 'chmod 700 memv2.sh'")
print("\nDone.")
print("\n")
print("\nRun memv2.sh (on ec2 instance)")
system(f"ssh -i key6.pem ec2-user@{ip_address} './memv2.sh'")
print("\nend of memv2 script")
print("\n")

# list files in instance
print("\nList files in instance" + new_ec2_instance[0].id + "...")
system(f"ssh -i key6.pem ec2-user@{ip_address} 'ls -l'")
print("\nDone.")
print("\n")

# # copy install-node.sh and run it
# print("\nCopying install-node.sh to ec2 instance")
# system(
#     f"scp -o StrictHostKeyChecking=no -i key6.pem install-node.sh ec2-
user@{ip_address}:".
# )
# print("\nSet permissions on install-node.sh")
# system(f"ssh -i key6.pem ec2-user@{ip_address} 'chmod 700 install-
node.sh'")

```



```

# print("\nRun install-node.sh")
# system(f"ssh -i key6.pem ec2-user@{ip_address} './install-node.sh'")
# print("\nend of install-node.sh script")
# print("\n")

# # install Node.js
# print("\nInstalling Node.js...")
# system("curl -o- https://raw.githubusercontent.com/nvm-
sh/nvm/v0.39.5/install.sh | bash")
# system("source ~/.nvm/nvm.sh")
# system("nvm install 16")
# system("node -v")
# print("\nDone.")

except Exception as e:
    print(e)

# Upload the app.js file to ec2 instance
file_path = "app.js"
print("\nUploading " + file_path + " to ec2 instance...")
system(
    f"scp -o StrictHostKeyChecking=no -i key6.pem {file_path} ec2-
user@{ip_address}:"
)
print("\nDone.")
print("now connect via SSH and install node, then run app")
time.sleep(10)

# # run app.js
# print("\nRunning app.js...")
# system("node app.js")
# print("\nDone.")

# # open browser window to ec2 instance
# print("\nOpening browser window to ec2 instance...")

# print("\nOpening webpage at: " + new_ec2_instance[0].public_ip_address)
# print("\n\n")
# ip_address = new_ec2_instance[
#     0
# ].public_ip_address # Set variable to instance public IP
# webbrowser.open(
#     "http://" + new_ec2_instance[0].public_ip_address + ":3000"
# ) # Open web browser to instance public IP

# # install Node.js

```

```

# print("\nInstalling Node.js...")
# system("curl -o- https://raw.githubusercontent.com/nvm-
sh/nvm/v0.39.5/install.sh | bash -")
# system("source ~/.nvm/nvm.sh")
# system("nvm install 16")
# #export nvm dir to bashrc
# system("export NVM_DIR=$HOME/.nvm")
# system("[ -s \"$NVM_DIR/nvm.sh\" ] && \. \"$NVM_DIR/nvm.sh\"")
# system("[ -s \"$NVM_DIR/bash_completion\" ] && \.
\"$NVM_DIR/bash_completion\"")
# system("nvm use 16")
# print("\ninstalling npm")
# system("npm install -g npm@latest")
# # get node version number
# print("\nNode version:")
# system("node -v")
# print("\nDone.")

# # Upload the app.js file
# file_path = "app.js"
# if os.path.exists(file_path):
#     try:
#         s3_client.upload_file(
#             file_path, bucket_name, "app.js", ExtraArgs={"ContentType":
"text/javascript"}
#         )
#         print(f"File {file_path} uploaded as app.js.")
#     except ClientError as e:
#         print(f"Error uploading {file_path}: {e}")
#         exit(1)
# else:
#     print(f"File {file_path} does not exist.")
#     exit(1)

# # run app.js
# print("\nRunning app.js...")
# system("node app.js")

#CloudWatch alarms setup

print("Setting up CloudWatch alarms...")
cloudwatch = boto3.client("cloudwatch")

# CPU utilization greater than 50%
try:
    cloudwatch.put_metric_alarm(

```

```

        AlarmName="HighCPUUtilization",
        ComparisonOperator="GreaterThanThreshold",
        EvaluationPeriods=1,
        MetricName="CPUUtilization",
        Namespace="AWS/EC2",
        Period=60,
        Statistic="Average",
        Threshold=50.0,
        ActionsEnabled=True,
        AlarmActions=[
            "arn:aws:automate:us-east-1:ec2:reboot",
        ],
        AlarmDescription="Alarm if server CPU utilization exceeds 50%",
        Dimensions=[
            {"Name": "InstanceId", "Value": new_ec2_instance[0].id},
        ],
        Unit="Percent",
    )
    print("High CPU utilization alarm created.")
except ClientError as e:
    print(f"Error creating high CPU utilization alarm: {e}")

# CPU utilization less than 30%
try:
    cloudwatch.put_metric_alarm(
        AlarmName="LowCPUUtilization",
        ComparisonOperator="LessThanThreshold",
        EvaluationPeriods=1,
        MetricName="CPUUtilization",
        Namespace="AWS/EC2",
        Period=60,
        Statistic="Average",
        Threshold=30.0,
        ActionsEnabled=True,
        AlarmActions=[
            "arn:aws:automate:us-east-1:ec2:terminate",
        ],
        AlarmDescription="Alarm if server CPU utilization below 30%",
        Dimensions=[
            {"Name": "InstanceId", "Value": new_ec2_instance[0].id},
        ],
        Unit="Percent",
    )
    print("Low CPU utilization alarm created.")
except ClientError as e:
    print(f"Error creating low CPU utilization alarm: {e}")

try:

```

```

# Describe alarms
print("\nDescribing alarms...")
response = cloudwatch.describe_alarms()
for alarm in response["MetricAlarms"]:
    print(f"Alarm Name: {alarm['AlarmName']}")
    print(f"Alarm Description: {alarm['AlarmDescription']}")
    print(f"Alarm State: {alarm['StateValue']}")
    print(f"Alarm Actions: {alarm['AlarmActions']}")
    print(f"Alarm Comparison: {alarm['ComparisonOperator']}")
    print(f"Evaluation Periods: {alarm['EvaluationPeriods']}")
    print(f"Metric Name: {alarm['MetricName']}")
    print(f"Namespace: {alarm['Namespace']}")
    print(f"Period: {alarm['Period']}")
    print(f"Statistic: {alarm['Statistic']}")
    print(f"Threshold: {alarm['Threshold']}")
    print(f"Unit: {alarm['Unit']}")
    print("\n")
except ClientError as e:
    print(f"Error describing alarms: {e}")

# print("\n\nexiting...")
# time.sleep(1)
exit()

```

## APPENDIX B: monitor.sh - Instance info.

```
#!/usr/bin/bash
# Michael MCKibbin 20092733
# Adapted from course provided script.
#
INSTANCE_ID=$(curl -s http://169.254.169.254/latest/meta-data/instance-id)
MEMORYUSAGE=$(free -m | awk 'NR==2{printf "%.2f%%", $3*100/$2 }')
PROCESSES=$(expr $(ps -A | grep -c .) - 1)
HTTPD_STATUS=$(systemctl status $1 | awk 'NR == 3')
HTTPD_UPTIME=$(uptime | awk '{print $3,$4}' | cut -d, -f1)
HTTPD_PROCESSES=$(ps -A | grep -c httpd)
APACHE_PROCESSES=$(ps -A | grep -c apache)
APACHE_PROCESSES_NoGREP=$(ps -ef | grep -v grep | grep -c apache)
HTTP_PORT=$(netstat -tuln | grep -c ":80")
HTTPS_PORT=$(netstat -tuln | grep -c ":443")
AVAILABILITY_ZONE=$(curl -s http://169.254.169.254/latest/meta-
data/placement/availability-zone)
AMI_ID=$(curl -s http://169.254.169.254/latest/meta-data/ami-id)
SECURITY_GROUP=$(curl -s http://169.254.169.254/latest/meta-data/security-
groups)
IPV4=$(curl -s http://169.254.169.254/latest/meta-data/public-ipv4)
SSH_PORT=$(netstat -tuln | grep -c ":22")

echo "Instance ID: $INSTANCE_ID"
echo "Memory utilisation: $MEMORYUSAGE"
if [ $HTTPD_PROCESSES -ge 1 ]
then
    echo "Web server is running"
else
    echo "Web server is NOT running"
fi

echo "No of HTTPD processes: $PROCESSES"
echo "HTTPD server status: $HTTPD_STATUS"
echo "HTTPD server uptime: $HTTPD_UPTIME"
echo "No of Apache processes: $APACHE_PROCESSES"
echo "No of Apache processes (No grep): $APACHE_PROCESSES_NoGREP"
echo "No of SSH connections: $SSH_PORT"
echo "No of HTTP connections: $HTTP_PORT"
echo "No of HTTPS connections: $HTTPS_PORT"
echo "Availability zone: $AVAILABILITY_ZONE"
echo "AMI ID: $AMI_ID"
echo "Security group: $SECURITY_GROUP"
echo "Public IPv4: $IPV4"
```

## APPENDIX C: memv2.sh - Cloudwatch.

```
#!/bin/bash
TOKEN=`curl -X PUT "http://169.254.169.254/latest/api/token" -H "X-aws-ec2-
metadata-token-ttl-seconds: 21600"`
INSTANCE_ID=$(curl -H "X-aws-ec2-metadata-token: $TOKEN"
http://169.254.169.254/latest/meta-data/instance-id)

#INSTANCE_ID=$(curl -s http://169.254.169.254/latest/meta-data/instance-id)
USEDMEMORY=$(free -m | awk 'NR==2{printf "%.2f\t", $3*100/$2 }')
TCP_CONN=$(netstat -an | wc -l)
TCP_CONN_PORT_80=$(netstat -an | grep 80 | wc -l)
IO_WAIT=$(iostat | awk 'NR==4 {print $5}')

aws cloudwatch put-metric-data --metric-name memory-usage --dimensions
Instance=$INSTANCE_ID --namespace "Custom" --value $USEDMEMORY
aws cloudwatch put-metric-data --metric-name Tcp_connections --dimensions
Instance=$INSTANCE_ID --namespace "Custom" --value $TCP_CONN
aws cloudwatch put-metric-data --metric-name TCP_connection_on_port_80 --
dimensions Instance=$INSTANCE_ID --namespace "Custom" --value
$TCP_CONN_PORT_80
aws cloudwatch put-metric-data --metric-name IO_WAIT --dimensions
Instance=$INSTANCE_ID --namespace "Custom" --value $IO_WAIT
```